

文件 编辑 视图 帮助

b3353@LEGIONR7000P: /mnt

nd in params_dict
Multi-thread loading shards: 100% Completed | 1/1 [00:01<00:00, 1.07s/it]
[2026-09-25 11:27:26] Load weight end. elapsed=1.94 s, type=Qwen3ForCausalLM, avail mem=5.65 GB, mem usage=1.18 GB.
[2026-09-25 11:27:29] KV Cache is allocated. dtype: torch.bfloat16, #tokens: 20914, K size: 1.12 GB, V size: 1.12 GB
[2026-09-25 11:27:29] Memory pool end. avail mem=3.04 GB
[2026-09-25 11:27:29] Capture target prefill CUDA graph begin. backend=tc_pieewise, num_tokens=[4, 8, 12, 16, 20, 24, 28, 32, 48, 64, 80, 96, 112, 128, 144, 160, 176, 192, 208, 224, 240, 256, 288, 320, 352, 384, 416, 448, 480, 512, 576, 640, 704, 768, 832, 896, 960, 1024, 1280, 1536, 1792, 2048], avail mem=2.47 GB
Compiling num tokens (num_tokens=2048): 0% | 0/42 [00:00<?, ?it/s]
2026-09-25 11:28:09] Compiling a graph for dynamic shape takes 0.27 s
[2026-09-25 11:28:12] Compiling a graph for dynamic shape takes 0.28 s
Compiling num tokens (num_tokens=4): 100% | 42/42 [00:07<00:00, 5.26it/s]
Capturing num tokens (num_tokens=4 avail_mem=1.63 GB): 100%
[2026-09-25 11:28:18] Capture target prefill CUDA graph begin. backend=tc_pieewise, num_tokens=[4, 8, 12, 16, 20, 24, 28, 32, 48, 64, 80, 96, 112, 128, 144, 160, 176, 192, 208, 224, 240, 256, 288, 320, 352, 384, 416, 448, 480, 512, 576, 640, 704, 768, 832, 896, 960, 1024, 1280, 1536, 1792, 2048], avail mem=1.58 GB
[2026-09-25 11:28:18] Capture target decode CUDA graph begin. backend=tc_pieewise, num_tokens=[4, 8, 12, 16, 20, 24, 28, 32, 48, 64, 80, 96, 112, 128, 144, 160, 176, 192, 208, 224, 240, 256, 288, 320, 352, 384, 416, 448, 480, 512, 576, 640, 704, 768, 832, 896, 960, 1024, 1280, 1536, 1792, 2048], avail mem=1.58 GB
Capturing batches (bs=1 avail_mem=1.50 GB): 100%
[2026-09-25 11:28:27] Capture target decode CUDA graph begin. backend=tc_pieewise, num_tokens=[4, 8, 12, 16, 20, 24, 28, 32, 48, 64, 80, 96, 112, 128, 144, 160, 176, 192, 208, 224, 240, 256, 288, 320, 352, 384, 416, 448, 480, 512, 576, 640, 704, 768, 832, 896, 960, 1024, 1280, 1536, 1792, 2048], avail mem=1.58 GB
[2026-09-25 11:28:27] max_total_num_tokens=20914, context_len=40960, available_gpu_mem=1.58 GB
[2026-09-25 11:28:27] Tree cache initialized: source=False streaming_wrapped=False
[2026-09-25 11:28:27] INFO: Started server process
[2026-09-25 11:28:27] INFO: Waiting for application
[2026-09-25 11:28:27] Using default chat sampling method 'top_p': 0.95
[2026-09-25 11:28:27] INFO: Application started
[2026-09-25 11:28:27] INFO: Uvicorn running on 127.0.0.1:48818
[2026-09-25 11:28:28] INFO: 127.0.0.1:48818 -
[2026-09-25 11:28:34] Prefill batch, #new-seq: 1,

介绍你自己。/no_think"

cd","object":"chat.completi

/Qwen3-0.6B","choices":

b3353@LEGIONR7000P: /mnt/c/Users/b3353\$ curl http://127.0.0.1:30000/v1/models

{ "object": "list", "data": [{ "id": "Qwen/Qwen3-0.6B", "object": "model", "created": 1790306936, "owned_by": "sglang", "root": "Qwen/Qwen3-0.6B", "parent": null, "max_model_len": 40960 }] }

b3353@LEGIONR7000P: /mnt/c/Users/b3353\$ curl http://127.0.0.1:30000/v1/chat/completions \

-H "Content-Type: application/json" \

-d '{

"model": "Qwen/Qwen3-0.6B",

"messages": [

{ "role": "user", "content": "请用一句话介绍你自己。/no_think" }

],

"max_tokens": 128

}'

^C

b3353@LEGIONR7000P: /mnt/c/Users/b3353\$

b3353@LEGIONR7000P: /mnt/c/Users/b3353\$ curl http://127.0.0.1:30000/v1/chat/completions \

-H "Content-Type: application/json" \

-d '{

"model": "Qwen/Qwen3-0.6B",

"messages": [

{ "role": "user", "content": "请用一句话介绍你自己。/no_think" }

],

"max_tokens": 128

}'

{ "id": "8cc05f9467604d7095081f00d8531dcd", "object": "chat.completion", "created": 1790307004, "model": "Qwen/Qwen3-0.6B", "choices": [{ "index": 0, "message": { "role": "assistant", "content": "<think>\n\n\n我是AI助手，旨在提供帮助和支持。", "reasoning_content": null, "tool_calls": null }, "logprobs": null, "finish_reason": "stop", "matched_stop": 151645 }], "usage": { "prompt": 128, "completion": 128, "total": 256 } }

b3353@LEGIONR7000P: /mnt/c/Users/b3353\$

随心输入

+ 帮我批准

GPT-6 Sol 轻度

最近

修5/5

你好

华为昇腾

安装 humanizer Agent Skill

W365

教学空间

workload_results.csv

+

文件

自动保存

开始

插入

页面

公式

数据

审阅

视图

工具

会员专享

WPS AI

替换

分享

格式刷

粘贴

宋体

11

A+

A-

B

I

U

ab

田

常规

转换

行和列

表格样式

条件格式

单元格样式

填充

求和

排序

筛选

冻结

查找

A1

fx

id

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W
1	id	trace_in	trace_out	target_in	target_out	input_tok	output_tok	status	latency_s														
2	1	40160	6	313	6	332	6	200	0.18														
3	2	36640	4	286	4	305	4	200	0.037														
4	3	23769	3	185	4																		
5	4	32140	92	251	48																		
6	5	18663	210	145	48																		
7	6	28	445	32	48																		
8	7	42608	7	320	7																		
9	8	23479	169	183	48																		
10	9	18053	16	141	16																		
11	10	16838	292	131	48																		
12	11	24	216	32	48																		
13	12	546	78	32	48																		

b3353@LEGIONR7000P: /mnt

(ray-env) b3353@LEGIONR7000P:/mnt/c/Users/b3353\$ head -n 2 mooncake_synthetic_trace.jsonl~

head: cannot open 'mooncake_synthetic_trace.jsonl~' for reading: No such file or directory

(ray-env) b3353@LEGIONR7000P:/mnt/c/Users/b3353\$ mv /mnt/c/Users/b3353/mooncake_synthetic_trace.jsonl /mnt/e/sglangtraining/

(ray-env) b3353@LEGIONR7000P:/mnt/c/Users/b3353\$ cd /mnt/e/sglangtraining

(ray-env) b3353@LEGIONR7000P:/mnt/e/sglangtraining\$ head -n 2 mooncake_synthetic_trace.jsonl

{"timestamp": 0, "input_length": 40160, "output_length": 6, "hash_ids": [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78]}

{"timestamp": 0, "input_length": 36640, "output_length": 4, "hash_ids": [79, 80, 81, 82, 83, 84, 85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95, 96, 97, 98, 99, 100, 101, 102, 103, 104, 105, 106, 107, 108, 109, 110, 111, 112, 113, 114, 115, 116, 117, 118, 119, 120, 121, 122, 123, 124, 125, 126, 127, 128, 129, 130, 131, 132, 133, 134, 135, 136, 137, 138, 139, 140, 141, 142, 143, 144, 145, 146, 147, 148, 149, 150]}

(ray-env) b3353@LEGIONR7000P:/mnt/e/sglangtraining\$ nano workload.py

(ray-env) b3353@LEGIONR7000P:/mnt/e/sglangtraining\$ python workload.py

id trace_in trace_out target_in target_out actual_in actual_out status latency_s

1 40160 6 313 6 332 6 200 0.18

2 36640 4 286 4 305 4 200 0.037

3 23769 3 185 4 204 4 200 0.032

4 32140 92 251 48 270 48 200 0.347

5 18663 210 145 48 164 48 200 0.318

6 28 445 32 48 51 48 200 0.275

7 42608 7 320 7 339 7 200 0.059

8 23479 169 183 48 202 48 200 0.275

9 18053 16 141 16 160 16 200 0.103

10 16838 292 131 48 150 48 200 0.251

11 24 216 32 48 51 48 200 0.298

12 546 78 32 48 51 48 200 0.45

(ray-env) b3353@LEGIONR7000P:/mnt/e/sglangtraining\$

workload_results

+

100%